

Adrian Cortez

LLM Engineering Architect | Agentic Systems | RAG | Evals & Reliability | Production AI

Manila, Philippines (Remote) • Open to Staff/Principal
LLM Engineering roles

adrian.cortez@samplemail.com • +63 900 000 0000 •
github.com/adriancortez (sample)

About

Staff-level LLM engineer focused on building reliable, production-grade AI systems. I design agentic workflows with tool calling, harden them with evals, tracing, and fallbacks, and improve factuality with retrieval and citation discipline. I optimize for measurable outcomes: quality, latency, cost, and operational robustness.

Focus: Agents, Tool Schemas, RAG, Reranking, Evals, Observability, Safety

Stack: Python (FastAPI, Pydantic, AsyncIO), TypeScript, Docker, Postgres/pgvector, Redis, OpenTelemetry

Experience

Staff LLM Engineer / Technical Lead | Stealth AI Studio (Product & Platform)

2023-06 - Present • Remote (APAC)

- Led the architecture for a tool-using agent platform (hub-and-spoke orchestration) with strict JSON contracts and action audit logs.
- Implemented reliability patterns: timeouts, retries with jitter, circuit breakers, and tiered fallback models to reduce failed runs by 65% (internal metric).
- Built an eval harness (golden sets + regression snapshots) for prompts, tools, and end-to-end tasks; shipped weekly quality gates for releases.
- Designed a RAG pipeline with hybrid retrieval + reranking + citation requirements; reduced unsupported claims in QA audits.
- Partnered with product/security on prompt-injection defenses and PII handling (redaction, retention, and access controls).

Senior Machine Learning Engineer (NLP) | CloudScale Labs

2020-02 - 2023-05 • Manila, Philippines

- Owned NLP services for summarization and extraction; migrated legacy pipelines to transformer-based architectures with measurable accuracy gains.
- Introduced structured outputs (schemas) and post-generation validators for downstream automation reliability.
- Built observability: trace IDs per request, token/cost accounting, and per-tool latency metrics; improved incident triage time.
- Productionized model serving with canary deploys and online evaluation hooks for A/B testing.

Software Engineer (Backend & Data) | SearchWorks (Information Retrieval Team)

2017-07 - 2020-01 • Manila, Philippines

- Built search and ranking features: indexing, query parsing, and relevance scoring with offline evaluation.

- Delivered a production-style ranking service with automated regression tests and data-quality checks.
- Collaborated with analytics to instrument funnels and establish performance baselines for releases.

Selected Projects

- **Agent Reliability Kit:** reusable policies for retries, fallbacks, circuit breakers, and structured audit logs; drop-in for FastAPI services.
- **RAG Quality Dashboard:** retrieval hit-rate, citation coverage, and hallucination flags; supports regression comparisons by build tag.
- **Schema-first Prompting:** prompt templates + JSON schema validation + repair loop for tool arguments and downstream workflows.

Education

MSc, Machine Learning (NLP/LLMs specialization)

Thesis: Reliability methods for tool-using LLM agents (schema constraints, verification steps, regression evals)

BSc, Computer Science

Focus: distributed systems, databases, algorithms, software engineering; capstone on search + ranking service with evals

Certifications

- MLOps Fundamentals (experiment tracking, serving, monitoring, data validation)
- Cloud Architecture (networking, IAM, observability, cost controls)

Top Skills

LLM Agents & Tool Calling

Evaluation & QA

Observability (Tracing/Metrics)

Python/FastAPI

Postgres/pgvector

RAG & Retrieval

Prompt + Schema Design

Production Reliability

TypeScript/Node

Security & PII Handling

Interests

- Mechanical keyboards and workflow ergonomics (layout experiments, macro tuning)
- Coffee brewing (pour-over/espresso) with repeatable experiments and logs
- Chess (rapid) and post-game analysis
- Long walks with technical audio learning
- Photography (street/architecture) and seasonal hiking